



## Credit Card Fraud Detection

In this project, the focus is on developing robust machine learning models to enhance the ability of credit card companies in recognizing fraudulent transactions, preventing unwarranted charges for customers. The dataset contains transformed data to maintain confidentiality, with untouched Time and Amount columns. The 'Class' feature, representing fraud (1) or legit (0), serves as the dependent variable. Our objective is to construct advanced models that effectively identify fraudulent activities, ensuring a secure financial environment for credit card holders.

### Reading the dataset

```
In [1]: 1 # Importing necessary Libraries
        2 import pandas as pd
        3 import numpy as np
        4 from sklearn.model_selection import train_test_split
        5 from sklearn.linear_model import LogisticRegression
        6 from sklearn.tree import DecisionTreeClassifier
        7 from sklearn.ensemble import RandomForestClassifier
        8 from sklearn.metrics import accuracy_score
```

```
In [2]: 1 # Importing data file
        2 data = pd.read_csv('creditcard.csv')
        3 data.head()
```

```
Out[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431

5 rows × 31 columns



In [3]: 1 data.tail()

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...
<b>284802</b>	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...
<b>284803</b>	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...
<b>284804</b>	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...
<b>284805</b>	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...
<b>284806</b>	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...

5 rows × 31 columns

In [4]: 1 *# Information about the dataset*  
2 data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Time    284807 non-null    float64
1    V1       284807 non-null    float64
2    V2       284807 non-null    float64
3    V3       284807 non-null    float64
4    V4       284807 non-null    float64
5    V5       284807 non-null    float64
6    V6       284807 non-null    float64
7    V7       284807 non-null    float64
8    V8       284807 non-null    float64
9    V9       284807 non-null    float64
10   V10      284807 non-null    float64
11   V11      284807 non-null    float64
12   V12      284807 non-null    float64
13   V13      284807 non-null    float64
14   V14      284807 non-null    float64
15   V15      284807 non-null    float64
16   V16      284807 non-null    float64
17   V17      284807 non-null    float64
18   V18      284807 non-null    float64
19   V19      284807 non-null    float64
20   V20      284807 non-null    float64
21   V21      284807 non-null    float64
22   V22      284807 non-null    float64
23   V23      284807 non-null    float64
24   V24      284807 non-null    float64
25   V25      284807 non-null    float64
26   V26      284807 non-null    float64
27   V27      284807 non-null    float64
28   V28      284807 non-null    float64
29   Amount   284807 non-null    float64
30   Class    284807 non-null    int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [5]: 1 # Checking null values
        2 data.isnull().sum()
```

```
Out[5]: Time      0
        V1        0
        V2        0
        V3        0
        V4        0
        V5        0
        V6        0
        V7        0
        V8        0
        V9        0
        V10       0
        V11       0
        V12       0
        V13       0
        V14       0
        V15       0
        V16       0
        V17       0
        V18       0
        V19       0
        V20       0
        V21       0
        V22       0
        V23       0
        V24       0
        V25       0
        V26       0
        V27       0
        V28       0
        Amount    0
        Class     0
        dtype: int64
```

## Splitting the data into 'Legit' and 'Fraudulent' transactions

```
In [6]: 1 # Distribution of legit and fraudulent transactions
        2 data['Class'].value_counts()
```

```
Out[6]: Class
0      284315
1        492
Name: count, dtype: int64
```

This is a highly unbalanced dataset as more than 99% of the data has legit transactions and we can not feed this data to the machine learning model. Hence, we will make some necessary changes in the data.

**0 = Legit Transaction , 1 = Fraudulent Transaction**

```
In [7]: 1 # Seperating data for analysis
        2 legit = data[data.Class == 0]
        3 fraud = data[data.Class == 1]
```

```
In [8]: 1 print(legit.shape)
        2 print(fraud.shape)
```

```
(284315, 31)
(492, 31)
```

```
In [9]: 1 # Statistical measures of data
        2
        3 legit.Amount.describe()
```

```
Out[9]: count    284315.000000
        mean      88.291022
        std       250.105092
        min        0.000000
        25%        5.650000
        50%       22.000000
        75%       77.050000
        max     25691.160000
        Name: Amount, dtype: float64
```

```
In [10]: 1 fraud.Amount.describe()
```

```
Out[10]: count    492.000000
         mean    122.211321
         std     256.683288
         min      0.000000
         25%      1.000000
         50%      9.250000
         75%    105.890000
         max    2125.870000
         Name: Amount, dtype: float64
```

```
In [11]: 1 # Compare the values for both transactions
        2
        3 data.groupby('Class').mean()
```

```
Out[11]:
```

	V8	V9 ...	V20	V21	V22	V23	V24	V25	V26	V27	V28	
00987	0.004467	...	-0.000644	-0.001235	-0.000024	0.000070	0.000182	-0.000072	-0.000089	-0.000295	-0.000131	8
570636	-2.581123	...	0.372319	0.713588	0.014049	-0.040308	-0.105130	0.041449	0.051648	0.170575	0.075667	12

## Undersampling

Reducing the number of legit transactions to fraudulent transactions to build a robust machine learning model i.e. 492 transactions.

```
In [48]: 1 legit_sample = legit.sample(n=492)
         2 legit_sample
```

Out[48]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...
<b>32508</b>	36831.0	-0.409064	-0.357429	1.996208	-0.986927	-0.231251	1.792335	-1.045989	0.807461	-1.267367	...
<b>263764</b>	161106.0	2.182664	-0.816907	-1.182612	-0.506839	-0.560771	-0.470415	-0.621406	-0.101797	-0.098331	...
<b>150439</b>	93353.0	-0.834929	0.195096	2.034287	-2.298311	-0.322921	0.019517	0.068651	-0.059660	-0.023583	...
<b>26708</b>	34215.0	1.216135	-0.165286	-0.080163	-0.855132	-0.346520	-0.736451	0.065392	-0.046705	0.968275	...
<b>45557</b>	42400.0	1.078540	-0.562184	1.199141	1.297798	-0.865169	1.374870	-1.262026	0.740642	1.564910	...
...	...	...	...	...	...	...	...	...	...	...	...
<b>231524</b>	146796.0	-3.316907	-3.998503	1.054058	-0.749884	1.240904	-1.371696	-2.249025	0.934289	-0.095539	...
<b>106572</b>	70019.0	1.356407	-1.324575	0.730767	-1.255086	-1.929883	-0.779525	-1.169145	-0.148030	-1.573841	...
<b>71575</b>	54347.0	0.550432	-0.849224	-0.028740	1.600370	-0.750556	-0.665107	0.531738	-0.107686	0.342753	...
<b>17191</b>	28512.0	-0.578161	1.604349	0.051961	0.594239	0.622138	-0.250479	0.449709	0.454511	-1.380644	...
<b>19082</b>	30003.0	-1.342081	1.000854	1.875593	-0.774800	0.347329	0.377771	0.756330	0.140343	0.280663	...

492 rows × 31 columns

```
In [49]: 1 # Creating a new dataset by merging sample legit data and fraudulent data
         2
         3 new_data = pd.concat([legit_sample, fraud], axis=0)
         4 new_data
```

Out[49]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...
<b>32508</b>	36831.0	-0.409064	-0.357429	1.996208	-0.986927	-0.231251	1.792335	-1.045989	0.807461	-1.267367	...
<b>263764</b>	161106.0	2.182664	-0.816907	-1.182612	-0.506839	-0.560771	-0.470415	-0.621406	-0.101797	-0.098331	...
<b>150439</b>	93353.0	-0.834929	0.195096	2.034287	-2.298311	-0.322921	0.019517	0.068651	-0.059660	-0.023583	...
<b>26708</b>	34215.0	1.216135	-0.165286	-0.080163	-0.855132	-0.346520	-0.736451	0.065392	-0.046705	0.968275	...
<b>45557</b>	42400.0	1.078540	-0.562184	1.199141	1.297798	-0.865169	1.374870	-1.262026	0.740642	1.564910	...
...	...	...	...	...	...	...	...	...	...	...	...
<b>279863</b>	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850	0.697211	-2.064945	...
<b>280143</b>	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170	0.248525	-1.127396	...
<b>280149</b>	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739	1.210158	-0.652250	...
<b>281144</b>	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002	1.058733	-1.632333	...
<b>281674</b>	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050	-0.068384	0.577829	...

984 rows × 31 columns

```
In [50]: 1 # Uniformly distributed the data
         2 new_data['Class'].value_counts()
```

Out[50]: Class  
0 492  
1 492  
Name: count, dtype: int64

```
In [51]: 1 new_data.groupby('Class').mean()
```

```
Out[51]:
```

	V8	V9	...	V20	V21	V22	V23	V24	V25	V26	V27	V28
0.044744	0.041445	...	0.016010	-0.057506	-0.063466	-0.029693	0.026263	-0.005472	-0.005357	0.037518	-0.007669	8
0.570636	-2.581123	...	0.372319	0.713588	0.014049	-0.040308	-0.105130	0.041449	0.051648	0.170575	0.075667	12

After grouping the new data on the basis of 'Class' we are getting a mean of USD 86.658 for legit transactions and in the original dataset it was USD 88.291. Hence, we can conclude that features of the data has not changed much.

## Splitting the data into dependnt and independent variables

```
In [52]: 1 x = new_data.drop('Class', axis=1)
2 y = new_data['Class']
```

```
In [53]: 1 print(x)
```

	Time	V1	V2	V3	V4	V5	V6	\
32508	36831.0	-0.409064	-0.357429	1.996208	-0.986927	-0.231251	1.792335	
263764	161106.0	2.182664	-0.816907	-1.182612	-0.506839	-0.560771	-0.470415	
150439	93353.0	-0.834929	0.195096	2.034287	-2.298311	-0.322921	0.019517	
26708	34215.0	1.216135	-0.165286	-0.080163	-0.855132	-0.346520	-0.736451	
45557	42400.0	1.078540	-0.562184	1.199141	1.297798	-0.865169	1.374870	
...	...	...	...	...	...	...	...	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	
...	...	...	...	...	...	...	...	
32508	-1.045989	0.807461	-1.267367	...	0.133482	0.669782	1.949540	
263764	-0.621406	-0.101797	-0.098331	...	-0.647472	-0.429267	-0.711759	
150439	0.068651	-0.059660	-0.023583	...	0.088986	0.298572	1.000051	
26708	0.065392	-0.046705	0.968275	...	-0.201874	0.068298	0.440049	
45557	-1.262026	0.740642	1.564910	...	-0.337736	-0.016064	0.142534	
...	...	...	...	...	...	...	...	
279863	-0.882850	0.697211	-2.064945	...	1.252967	0.778584	-0.319189	
280143	-1.413170	0.248525	-1.127396	...	0.226138	0.370612	0.028234	
280149	-2.234739	1.210158	-0.652250	...	0.247968	0.751826	0.834108	
281144	-2.208002	1.058733	-1.632333	...	0.306271	0.583276	-0.269209	
281674	0.223050	-0.068384	0.577829	...	-0.017652	-0.164350	-0.295135	
...	...	...	...	...	...	...	...	
32508	-0.044076	-0.988546	-0.751330	0.136308	0.242389	0.128163	11.00	
263764	0.299587	0.436893	-0.297904	0.544021	-0.056122	-0.053982	10.00	
150439	-0.501885	-0.329591	0.704233	-0.079623	-0.114027	0.026307	51.75	
26708	-0.146121	0.252030	0.815678	-0.548571	0.046713	-0.001224	1.00	
45557	-0.118760	-0.870898	0.382507	-0.194808	0.083445	0.011525	13.35	
...	...	...	...	...	...	...	...	
279863	0.639419	-0.294885	0.537503	0.788395	0.292680	0.147968	390.00	
280143	-0.145640	-0.081049	0.521875	0.739467	0.389152	0.186637	0.76	
280149	0.190944	0.032070	-0.739695	0.471111	0.385107	0.194361	77.89	
281144	-0.456108	-0.183659	-0.328168	0.606116	0.884876	-0.253700	245.00	
281674	-0.072173	-0.450261	0.313267	-0.289617	0.002988	-0.015309	42.53	

[984 rows x 30 columns]

In [54]: 1 `print(y)`

```
32508      0
263764     0
150439     0
26708      0
45557      0
..
279863     1
280143     1
280149     1
281144     1
281674     1
Name: Class, Length: 984, dtype: int64
```

In [55]: 1 `# Split the data onto train and test data`  
2  
3 `x_train, x_test, y_train, y_test = train_test_split(x, y, train_size= 0.2, stratify= y, ran`

In [56]: 1 `print(x.shape, x_train.shape, x_test.shape)`

```
(984, 30) (196, 30) (788, 30)
```

## Model training and evaluation

### Logistic Regression

In [57]: 1 `LR = LogisticRegression()`

In [58]: 1 `# Training the Logistic regression model with training data`  
2 `LR.fit(x_train, y_train)`

Out[58]: `LogisticRegression`  
`LogisticRegression()`

In [59]: 1 `# Accuracy on training data`  
2  
3 `x_train_prediction = LR.predict(x_train)`  
4 `training_data_accuracy = accuracy_score(x_train_prediction, y_train)`

In [60]: 1 `print("Accuracy score on training data using Logistic Regression :",training_data_accuracy)`  
  
Accuracy score on training data using Logistic Regression : 0.9285714285714286

In [61]: 1 `# Accuracy on test data`  
2  
3 `x_test_prediction = LR.predict(x_test)`  
4 `test_data_accuracy = accuracy_score(x_test_prediction, y_test)`

In [62]: 1 `print("Accuracy score on test data using Logistic Regression :",test_data_accuracy )`  
  
Accuracy score on test data using Logistic Regression : 0.9137055837563451

### Random Forrest Classifier

In [63]: 1 `RFC = RandomForestClassifier()`

```
In [64]: 1 RFC.fit(x_train, y_train)
```

```
Out[64]: ▾ RandomForestClassifier  
RandomForestClassifier()
```

```
In [65]: 1 x_train_prediction = RFC.predict(x_train)  
2 train_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
In [66]: 1 print("Accuracy score on train data using Random Forrest Classifier :",train_data_accuracy)  
  
Accuracy score on train data using Random Forrest Classifier : 1.0
```

```
In [67]: 1 x_test_prediction = RFC.predict(x_test)  
2 test_data_accuracy = accuracy_score(x_test_prediction, y_test)
```

```
In [68]: 1 print("Accuracy score on test data using Random Forrest Classifier :",test_data_accuracy )  
  
Accuracy score on test data using Random Forrest Classifier : 0.9276649746192893
```

## Decision Tree Classifier

```
In [38]: 1 DTC = DecisionTreeClassifier()
```

```
In [39]: 1 DTC.fit(x_train, y_train)
```

```
Out[39]: ▾ DecisionTreeClassifier  
DecisionTreeClassifier()
```

```
In [69]: 1 x_train_prediction = DTC.predict(x_train)  
2 train_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
In [70]: 1 print("Accuracy score on train data using Decision Tree Classifier :",train_data_accuracy )  
  
Accuracy score on train data using Decision Tree Classifier : 0.9540816326530612
```

```
In [71]: 1 x_test_prediction = DTC.predict(x_test)  
2 test_data_accuracy = accuracy_score(x_test_prediction, y_test)
```

```
In [72]: 1 print("Accuracy score on test data using Random Forrest Classifier :",test_data_accuracy )  
  
Accuracy score on test data using Random Forrest Classifier : 0.8946700507614214
```

**From the above model testing we can evaluate that 'Random Forest Classifier' model is the most accurate with accuracy of 92.76 %. So, we will use this model to predict future credit card fraud detections.**

## Conclusion

In this project, we classified transactions into legitimate and fraudulent categories, employing undersampling to balance the dataset and construct a resilient machine learning model. Logistic Regression, Random Forest Classifier, and Decision Tree Classifier models were implemented and evaluated. Among them, the 'Random Forest Classifier' emerged as the most accurate, boasting an impressive 92.76% accuracy. Consequently, we have chosen this model as our primary tool for predicting and detecting future instances of credit card fraud, offering a robust solution for enhancing security and protecting users from unauthorized transactions.